

How to Analyze and Tune MySQL Queries for Better Performance

Øystein Grøvlen
Senior Principal Software Engineer
MySQL Optimizer Team, Oracle
February 25, 2015

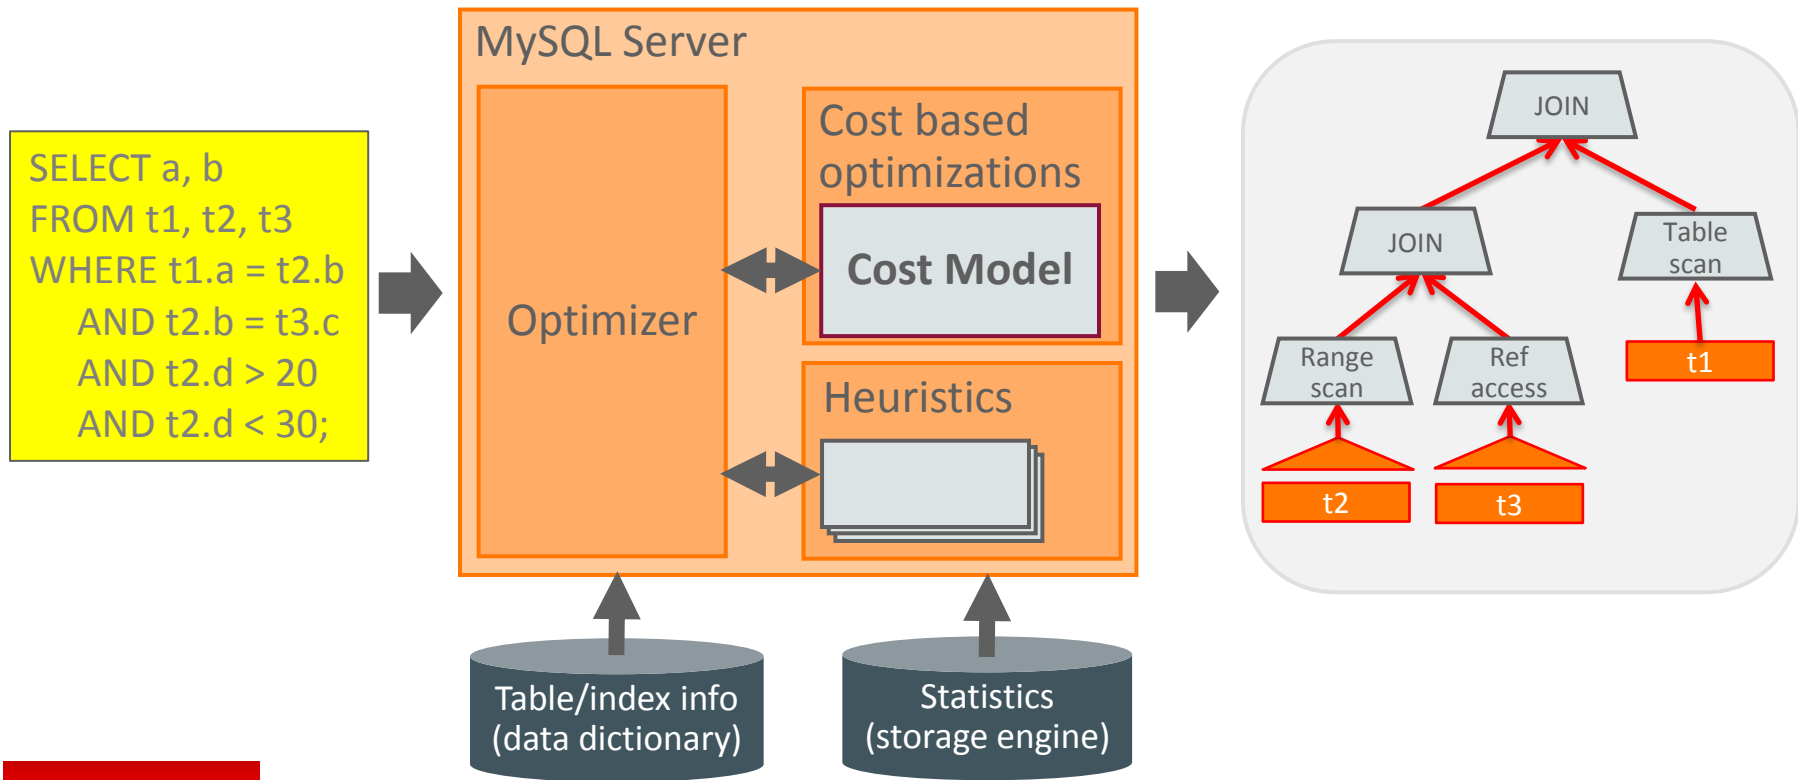
Program Agenda

- 1 Introduction to MySQL cost-based optimizer
- 2 Selecting data access method
- 3 Join optimizer
- 4 Sorting
- 5 Tools for monitoring, analyzing, and tuning queries
- 6 Influencing the Optimizer

Program Agenda

- 1 Introduction to MySQL optimizer
- 2 Selecting data access method
- 3 Join optimizer
- 4 Sorting
- 5 Tools for monitoring, analyzing, and tuning queries
- 6 Influencing the Optimizer

MySQL Optimizer



Cost-based Query Optimization


General idea

- Assign cost to operations
- Computes cost of partial or alternative plans
- Search for plan with lowest cost
- Cost-based optimizations:

Access method

Join order

Subquery strategy



Input to Cost Model

- **IO-cost:**

- Estimates from storage engine based on number of pages to read
- Both index and data pages

- **Schema:**

- Length of records and keys
- Uniqueness for indexes
- Nullability

- **Statistics:**

- Number of rows in table
- Key distribution/Cardinality:
 - Average number of records per key value
 - Only for indexed columns
 - Maintained by storage engine
- Number of records in an index range

Cost Model Example

```
SELECT SUM(o_totalprice) FROM orders
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';
```


Table scan:

- IO-cost: #pages in table
- CPU cost: #rows * ROW_EVALUATE_COST



Range scan (on secondary index):

- IO-cost: #pages to read from index + #rows_in_range
- CPU cost: #rows_in_range * ROW_EVALUATE_COST



Cost Model

Example

EXPLAIN SELECT SUM(o_totalprice) FROM orders
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|--------|------|---------------|------|---------|------|----------|-------------|
| 1 | SIMPLE | orders | ALL | i_o_orderdate | NULL | NULL | NULL | 15000000 | Using where |

EXPLAIN SELECT SUM(o_totalprice) FROM orders
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-06-30';

| Id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|--------|-------|---------------|---------------|---------|------|---------|-----------------------|
| 1 | SIMPLE | orders | range | i_o_orderdate | i_o_orderdate | 4 | NULL | 2235118 | Using index condition |

Cost Model Example: Optimizer Trace

join_optimization / row_estimation / table : orders / range_analysis

```
"table_scan": {
  "rows": 15000000,
  "cost": 3.12e6
} /* table_scan */,
"potential_range_indices": [
  {
    "index": "PRIMARY",
    "usable": false,
    "cause": "not_applicable"
  },
  {
    "index": "i_o_orderdate",
    "usable": true,
    "key_parts": [ "o_orderDATE", "o_orderkey" ]
  }
] /* potential_range_indices */,
...
```

```
"analyzing_range_alternatives": {
  "range_scan_alternatives": [
    {
      "index": "i_o_orderdate",
      "ranges": [ "1994-01-01 <= o_orderDATE <= 1994-12-31"
    ],
      "index_dives_for_eq_ranges": true,
      "rowid_ordered": false,
      "using_mrr": false,
      "index_only": false,
      "rows": 4489990,
      "cost": 5.39e6,
      "chosen": false,
      "cause": "cost"
    }
  ] /* range_scan_alternatives */,
  ...
} /* analyzing_range_alternatives */
```

Cost Model vs Real World

Measured Execution Times

| | Data in Memory | Data on Disk | Data on SSD |
|------------|--------------------|--------------|-------------|
| Table scan | 6.8 seconds | 36 seconds | 15 seconds |
| Index scan | 5.2 seconds | 2.5 hours | 30 minutes |

Force Index Scan:

SELECT SUM(o_totalprice)

FROM orders FORCE INDEX (i_o_orderdate)

WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';

Performance Schema

Disk I/O

```
SELECT event_name, count_read, avg_timer_read/1000000000.0 "Avg Read Time (ms)",
       sum_number_of_bytes_read "Bytes Read"
FROM performance_schema.file_summary_by_event_name
WHERE event_name='wait/io/file/innodb/innodb_data_file';
```

Table Scan

| event_name | count_read | Avg Read Time (ms) | Bytes Read |
|--------------------------------------|------------|--------------------|------------|
| wait/io/file/innodb/innodb_data_file | 115769 | 0.0342 | 1896759296 |

Index Scan

| event_name | count_read | Avg Read Time (ms) | Bytes Read |
|--------------------------------------|------------|--------------------|-------------|
| wait/io/file/innodb/innodb_data_file | 2188853 | 4.2094 | 35862167552 |

Program Agenda

- 1 Introduction to MySQL optimizer
- 2 Selecting data access method
- 3 Join optimizer
- 4 Sorting
- 5 Tools for monitoring, analyzing, and tuning queries
- 6 Influencing the Optimizer

Selecting Access Method

Finding the optimal method to read data from storage engine

- For each table, find the best access method:
 - Check if the access method is useful
 - Estimate cost of using access method
 - Select the cheapest to be used
- Choice of access method is cost based

Main access methods:

- Table scan
- Index scan
- Ref access
- Range scan
- Index merge
- Loose index scan

Ref Access

Single Table Queries

EXPLAIN SELECT * FROM customer WHERE c_custkey = 570887;

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|----------|-------|---------------|---------|---------|-------|------|-------|
| 1 | SIMPLE | customer | const | PRIMARY | PRIMARY | 4 | const | 1 | NULL |

EXPLAIN SELECT * FROM orders WHERE o_orderdate = '1992-09-12';

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|--------|------|---------------|---------------|---------|-------|------|-------|
| 1 | SIMPLE | orders | ref | i_o_orderdate | i_o_orderdate | 4 | const | 6271 | NULL |

Ref Access

Join Queries

EXPLAIN SELECT *

FROM orders JOIN customer ON c_custkey = o_custkey

WHERE o_orderdate = '1992-09-12';

| Id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|----------|--------|-------------------------------|---------------|---------|---------------------------|------|-------------|
| 1 | SIMPLE | orders | ref | i_o_orderdate, i_o_custkey | i_o_orderdate | 4 | const | 6271 | Using where |
| 1 | SIMPLE | customer | eq_ref | PRIMARY | PRIMARY | 4 | dbt3.orders. o_custkey | 1 | NULL |

Ref Access

Join Queries, continued

EXPLAIN SELECT *

FROM orders JOIN customer ON c_custkey = o_custkey

WHERE c_acctbal < -1000;

| Id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|----------|------|---------------|-------------|---------|-------------------------|---------|-------------|
| 1 | SIMPLE | customer | ALL | PRIMARY | NULL | NULL | NULL | 1500000 | Using where |
| 1 | SIMPLE | orders | ref | i_o_custkey | i_o_custkey | 5 | dbt3.customer.c_custkey | 7 | NULL |

Range Optimizer

- Goal: find the “minimal” ranges for each index that needs to be read
- Example:

SELECT * FROM t1 WHERE (key1 > 10 AND key1 < 20) AND key2 > 30

- Range scan using INDEX(key1):



- Range scan using INDEX(key2):



Range Optimizer

Optimizer Trace show ranges

```
SELECT a, b FROM t1
WHERE a > 10
      AND a < 25
      AND a NOT IN (11, 19))
      AND (b < 5 OR b > 10);
```

```
"analyzing_range_alternatives": {
  "range_scan_alternatives": [
    {
      "index": "i_a",
      "ranges": [
        "10 < a < 11",
        "11 < a < 19",
        "19 < a < 25"
      ],
      "index_dives_for_eq_ranges": true,
      "rowid_ordered": false,
      "using_mrr": false,
      "index_only": false,
      "rows": 3,
      "cost": 6.61,
      "chosen": true
    },
    {
      "index": "i_b",
      "ranges": [
        "NULL < b < 5",
        "10 < b"
      ],
      "index_dives_for_eq_ranges": true,
      "rowid_ordered": false,
      ...
    }
  ]
}
```

Range Optimizer: Case Study

Why table scan?

```
SELECT * FROM orders
WHERE YEAR(o_orderdate) = 1997 AND MONTH(o_orderdate) = 5
      AND o_clerk LIKE '%01866';
```

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|--------|------|---------------|------|---------|------|----------|-------------|
| 1 | SIMPLE | orders | ALL | NULL | NULL | NULL | NULL | 15000000 | Using where |

Index not considered

```
mysql> SELECT * FROM orders WHERE year(o_orderdate) = 1997 AND MONTH(...)
...
15 rows in set (8.91 sec)
```

Range Optimizer: Case Study

Rewrite query to avoid functions on indexed columns

```
SELECT * FROM orders
WHERE o_orderdate BETWEEN '1997-05-01' AND '1997-05-31'
AND o_clerk LIKE '%01866';
```

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|--------|-------|---------------|---------------|---------|------|--------|---------------------------------------|
| 1 | SIMPLE | orders | range | i_o_orderdate | i_o_orderdate | 4 | NULL | 376352 | Using index condition; Using where |

```
mysql> SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' AND ...
...
15 rows in set (0.91 sec)
```

Range Optimizer: Case Study

Adding another index

CREATE INDEX i_o_clerk ON orders(o_clerk);

SELECT * FROM orders

WHERE o_orderdate BETWEEN '1997-05-01' AND '1997-05-31'

AND o_clerk LIKE '%01866';

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|--------|-------|---------------|---------------|---------|------|--------|---------------------------------------|
| 1 | SIMPLE | orders | range | i_o_orderdate | i_o_orderdate | 4 | NULL | 376352 | Using index condition; Using where |

New index not considered

Range Optimizer: Case Study

Rewrite query, again

```
SELECT * FROM orders
WHERE o_orderdate BETWEEN '1997-05-01' AND '1997-05-31'
      AND o_clerk = 'Clerk#000001866';
```

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|--------|-------|-----------------------------|-----------|---------|------|------|---------------------------------------|
| 1 | SIMPLE | orders | range | i_o_orderdate, i_o_clerk | i_o_clerk | 16 | NULL | 1504 | Using index condition; Using where |

```
mysql> SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' AND ...
...
15 rows in set (0.01 sec)
```

Range Access for Multi-part Index

Example table with multi-part index

- Table:



- INDEX idx(a, b, c);



- Logical storage layout of index:



Range Optimizer: Case Study

Create multi-column index

```
CREATE INDEX i_o_clerk_date ON orders(o_clerk, o_orderdate);
```

```
SELECT * FROM orders  
WHERE o_orderdate BETWEEN '1997-05-01' AND '1997-05-31'  
AND o_clerk = 'Clerk#000001866';
```

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|--------|-------|--|----------------|---------|------|------|-----------------------|
| 1 | SIMPLE | orders | range | i_o_orderdate, i_o_clerk, i_o_clerk_date | i_o_clerk_date | 20 | NULL | 14 | Using index condition |

```
mysql> SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' AND ...  
...  
15 rows in set (0.00 sec)
```


Performance Schema: Query History

UPDATE performance_schema.setup_consumers

SET enabled='YES' WHERE name = 'events_statements_history';

```
mysql> SELECT sql_text, (timer_wait)/1000000000.0 "Time (ms)", rows_examined FROM
performance_schema.events_statements_history ORDER BY timer_start;
```

| sql_text | Time (ms) | rows_examined |
|---|------------|---------------|
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 8.1690 | 1505 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 7.2120 | 1505 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 8.1613 | 1505 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 7.0535 | 1505 |
| CREATE INDEX i_o_clerk_date ON orders(o_clerk,o_orderdate) | 82036.4190 | 0 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.7259 | 15 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.5791 | 15 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.5423 | 15 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.6031 | 15 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.2710 | 15 |


Program Agenda

- 1 Introduction to MySQL optimizer
- 2 Selecting data access method
- 3 Join optimizer**
- 4 Sorting
- 5 Tools for monitoring, analyzing, and tuning queries
- 6 Influencing the Optimizer

Join Optimizer

"Greedy search strategy"

- Goal: Given a JOIN of N tables, find the best JOIN ordering
- Strategy:
 - Start with all 1-table plans
 - Expand each plan with remaining tables
 - Depth-first
 - If "cost of partial plan" > "cost of best plan":
 - "prune" plan
 - Heuristic pruning:
 - Prune less promising partial plans
 - May in rare cases miss most optimal plan (turn off with **set optimizer_prune_level = 0**)




Complexity and Cost of Join Optimizer

Join of N tables: $N!$ possible plans to evaluate

Heuristics to reduce the number of plans to evaluate:

- Use **optimizer_search_depth** to limit the number of tables to consider
- Pre-sort tables on *size* and *key dependency* order (Improved in MySQL 5.6)
- When adding the next table to a partial plan, add all tables that it has an equality reference to (New in MySQL 5.6)



Join Optimizer: Case study

DBT-3 Query 8: National Market Share Query

```

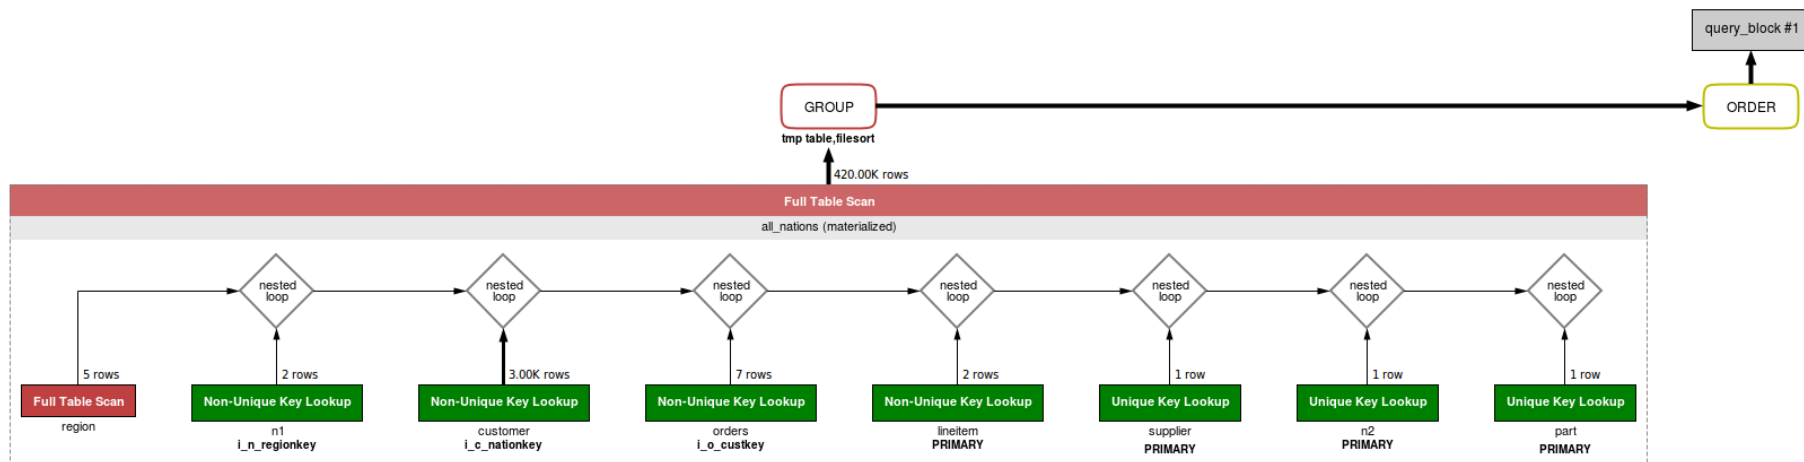
SELECT o_year, SUM(CASE WHEN nation = 'FRANCE' THEN volume ELSE 0 END) / SUM(volume) AS
    mkt_share
FROM (
    SELECT EXTRACT(YEAR FROM o_orderdate) AS o_year,
           l_extendedprice * (1 - l_discount) AS volume, n2.n_name AS nation
    FROM part
    JOIN lineitem ON p_partkey = l_partkey
    JOIN supplier ON s_suppkey = l_suppkey
    JOIN orders ON l_orderkey = o_orderkey
    JOIN customer ON o_custkey = c_custkey
    JOIN nation n1 ON c_nationkey = n1.n_nationkey
    JOIN region ON n1.n_regionkey = r_regionkey
    JOIN nation n2 ON s_nationkey = n2.n_nationkey
    WHERE r_name = 'EUROPE' AND o_orderdate BETWEEN '1995-01-01' AND '1996-12-31'
           AND p_type = 'PROMO BRUSHED STEEL'
    ) AS all_nations GROUP BY o_year ORDER BY o_year;

```

Join Optimizer: Case Study

MySQL Workbench: Visual EXPLAIN

Execution time: 3 min. 28 sec.



Join Optimizer: Case study

Force early processing of high selectivity predicates

```
SELECT o_year, SUM(CASE WHEN nation = 'FRANCE' THEN volume ELSE 0 END) / SUM(volume) AS
mkt_share
```

```
FROM (
```

```
SELECT EXTRACT(YEAR FROM o_orderdate) AS o_year,
l_extendedprice * (1 - l_discount) AS volume, n2.n_name AS nation
FROM part
```

part before lineitem

```
STRAIGHT_JOIN lineitem ON p_partkey = l_partkey
JOIN supplier ON s_suppkey = l_suppkey
JOIN orders ON l_orderkey = o_orderkey
JOIN customer ON o_custkey = c_custkey
JOIN nation n1 ON c_nationkey = n1.n_nationkey
JOIN region ON n1.n_regionkey = r_regionkey
JOIN nation n2 ON s_nationkey = n2.n_nationkey
WHERE r_name = 'EUROPE' AND o_orderdate BETWEEN '1995-01-01' AND '1996-12-31'
AND p_type = 'PROMO BRUSHED STEEL'
```

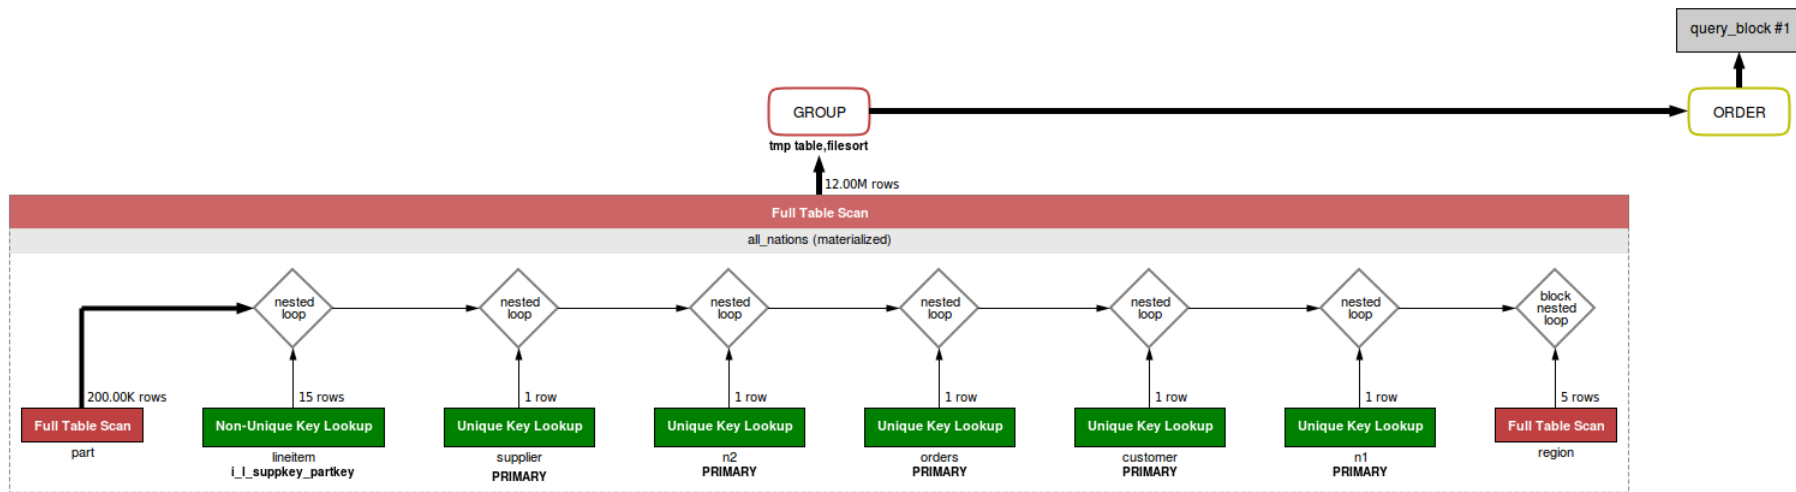
```
) AS all_nations GROUP BY o_year ORDER BY o_year;
```

Highest selectivity

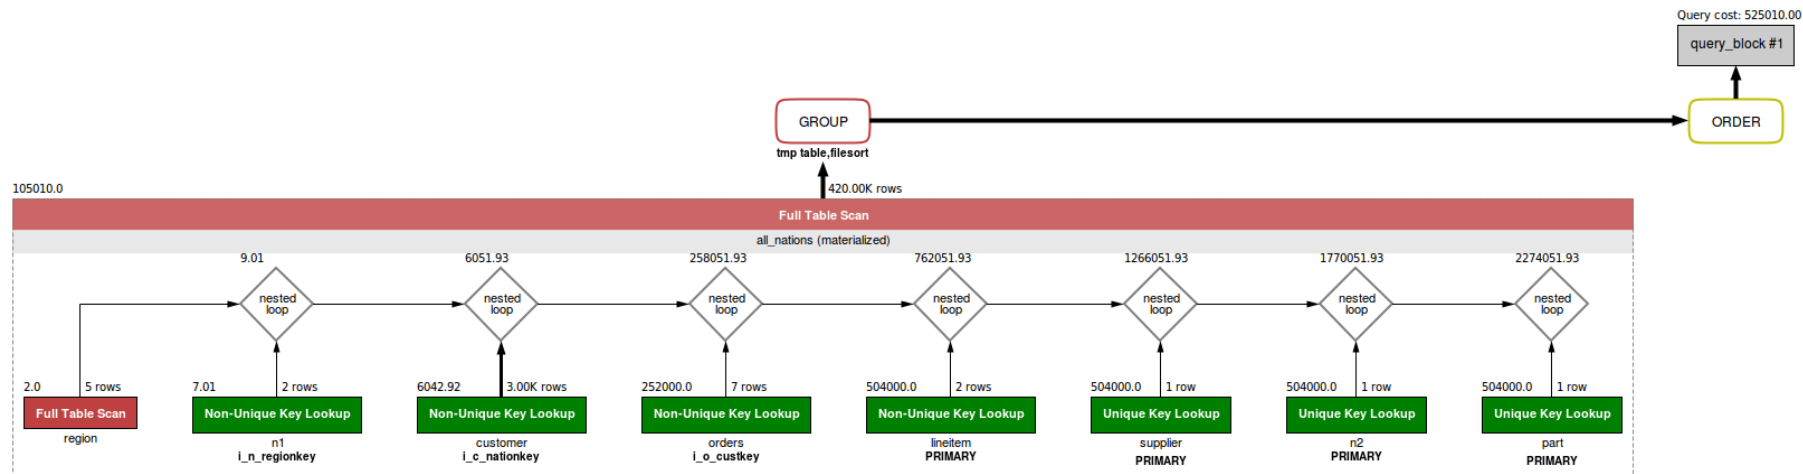
Join Optimizer: Case study

Improved join order

Execution time: 7 seconds



MySQL 5.7: Cost Information in Structured EXPLAIN



Program Agenda

- 1 Introduction to MySQL optimizer
- 2 Selecting data access method
- 3 Join optimizer
- 4 Sorting**
- 5 Tools for monitoring, analyzing, and tuning queries
- 6 Influencing the Optimizer

ORDER BY Optimizations

- General solution; “Filesort”:
 - Store query result in temporary table before sorting
 - If data volume is large, may need to sort in several passes with intermediate storage on disk.
- Optimizations:
 - Take advantage of index to generate query result in sorted order
 - For “LIMIT n ” queries, maintain priority queue of n top items in memory instead of filesort. (New in MySQL 5.6)

Filesort

SELECT * FROM orders ORDER BY o_totalprice ;

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|--------|------|---------------|------|---------|------|----------|----------------|
| 1 | SIMPLE | orders | ALL | NULL | NULL | NULL | NULL | 15000000 | Using filesort |

**SELECT c_name, o_orderkey, o_totalprice
FROM orders JOIN customer ON c_custkey = o_custkey
WHERE c_acctbal < -1000 ORDER BY o_totalprice ;**

| id | select type | table | type | possible keys | key | key len | ref | rows | extra |
|----|-------------|----------|------|---------------|-------------|---------|------|---------|--|
| 1 | SIMPLE | customer | ALL | PRIMARY | NULL | NULL | NULL | 1500000 | Using where; Using temporary; Using filesort |
| 1 | SIMPLE | orders | ref | i_o_custkey | i_o_custkey | 5 | ... | 7 | NULL |

Filesort

Status variables

Status variables related to sorting:

```
mysql> show status like 'Sort%';
```

| Variable_name | Value |
|-------------------|--------|
| Sort_merge_passes | 1 |
| Sort_range | 0 |
| Sort_rows | 136170 |
| Sort_scan | 1 |

>0: Intermediate storage on disk.

Consider increasing `sort_buffer_size`

Number of sort operations

(range scan or table/index scans)

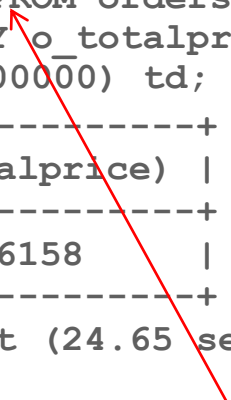
Number of rows sorted

Filesort: Case Study

```
mysql> FLUSH STATUS;


Query OK, 0 rows affected (0.00 sec)

mysql> SELECT AVG(o_totalprice) FROM (
  SELECT * FROM orders
    ORDER BY o_totalprice DESC
    LIMIT 100000) td;
+-----+
| AVG(o_totalprice) |
+-----+
| 398185.986158     |
+-----+
1 row in set (24.65 sec)
```



Unnecessary large data volume!

```
mysql> SHOW STATUS LIKE 'sort%';
+-----+-----+
| Variable_name | Value |
+-----+-----+
| Sort_merge_passes | 1432 |
| Sort_range       | 0    |
| Sort_rows        | 100000 |
| Sort_scan        | 1    |
+-----+-----+
4 rows in set (0.00 sec)
```



Many intermediate sorting steps!

Filesort: Case Study

Reduce amount of data to be sorted

```
mysql> FLUSH STATUS;
```

Query OK, 0 rows affected (0.00 sec)

```
mysql> SELECT AVG(o_totalprice) FROM (
    SELECT o_totalprice FROM orders
    ORDER BY o_totalprice DESC
    LIMIT 100000) td;
```

```
+-----+
| AVG(o_totalprice) |
+-----+
| 398185.986158      |
+-----+
```

1 row in set (8.18 sec)

```
mysql> SHOW STATUS LIKE 'sort%';
```

```
+-----+-----+
| Variable_name | Value |
+-----+-----+
| Sort_merge_passes | 229 |
| Sort_range       | 0 |
| Sort_rows        | 100000 |
| Sort_scan        | 1 |
+-----+-----+
```

4 rows in set (0.00 sec)

Filesort: Case Study

Increase sort buffer (1 MB)

Default is 256 kB



```
mysql> set sort_buffer_size=1024*1024;
```

```
mysql> FLUSH STATUS;
```

Query OK, 0 rows affected (0.00 sec)

```
mysql> SELECT AVG(o_totalprice) FROM
  ( SELECT o_totalprice FROM orders
    ORDER BY o_totalprice DESC
    LIMIT 100000) td;
```

| AVG(o_totalprice) |
|-------------------|
| 398185.986158 |

1 row in set (7.24 sec)

```
mysql> SHOW STATUS LIKE 'sort%';
```

| Variable_name | Value |
|-------------------|--------|
| Sort_merge_passes | 57 |
| Sort_range | 0 |
| Sort_rows | 100000 |
| Sort_scan | 1 |

4 rows in set (0.00 sec)

Filesort: Case Study

Increase sort buffer even more (8 MB)

```
mysql> set
      sort_buffer_size=8*1024*1024;

mysql> FLUSH STATUS;

Query OK, 0 rows affected (0.00 sec)

mysql> SELECT AVG(o_totalprice) FROM
      ( SELECT o_totalprice FROM orders
        ORDER BY o_totalprice DESC
        LIMIT 100000) td;
```

```
+-----+
| AVG(o_totalprice) |
+-----+
| 398185.986158      |
+-----+
1 row in set (6.30 sec)
```

```
mysql> SHOW STATUS LIKE 'sort%';

+-----+-----+
| Variable_name | Value |
+-----+-----+
| Sort_merge_passes |      0 |
| Sort_range      |      0 |
| Sort_rows       | 100000 |
| Sort_scan       |      1 |
+-----+-----+
4 rows in set (0.00 sec)
```

NB! Bigger sort buffer than needed will give unnecessary overhead

Use Index to Avoid Sorting

```
CREATE INDEX i_o_totalprice ON orders(o_totalprice);
```

```
SELECT AVG(o_totalprice) FROM
```

```
(SELECT o_totalprice FROM orders ORDER BY o_totalprice DESC LIMIT 100000) td;
```

| id | select type | table | Type | possible keys | key | key len | ref | rows | extra |
|----|-------------|------------|-------|---------------|----------------|---------|------|----------|-------------|
| 1 | PRIMARY | <derived2> | ALL | NULL | NULL | NULL | NULL | 100000 | NULL |
| 2 | DERIVED | orders | index | NULL | i_o_totalprice | 6 | NULL | 15000000 | Using index |

```
mysql> SELECT AVG(o_totalprice) FROM (
      SELECT o_totalprice FROM orders
      ORDER BY o_totalprice DESC LIMIT 100000) td;
```

```
...
```

```
1 row in set (0.06 sec)
```

Program Agenda

- 1 Introduction to MySQL optimizer
- 2 Selecting data access method
- 3 Join optimizer
- 4 Sorting
- 5 Tools for monitoring, analyzing, and tuning queries
- 6 Influencing the Optimizer

Useful tools

- MySQL Enterprise Monitor (MEM), Query Analyzer
 - Commercial product
- Performance schema, MySQL sys schema
- EXPLAIN
- EXPLAIN FORMAT=JSON
- Optimizer trace
- Slow log
- Status variables (SHOW STATUS LIKE 'Handler%')

MySQL Enterprise Monitor, Query Analyzer



ORACLE MySQL Enterprise Monitor

1 2 0 0 1 admin Refresh: Every Minute

Dashboards Events Query Analyzer Reports & Graphs Configuration

Browse Queries

Show 25 entries

Export data options...

Showing 1 to 25 of 1,197 entries

First

Previous

1

2

3

4

5

Next

Last

| Query | Database | | Counts | | | QRTI | Latency (hh:mm:ss.ms) | | | | | Rows | |
|--|----------|--|--------|-----|------|------|-----------------------|--------|-------|-------|-------------|--------|----------|
| | | | Exec | Err | Warn | | Total | Max | Avg | Locks | Avg History | Total | Examined |
| <div><div></div><div></div><div>COMMIT (1)</div></div> | mem | | 24,707 | 0 | 0 | 0.92 | 20:27.872 | 5.828 | 0.050 | 0.000 | <div></div> | 0 | 0 |
| <div><div></div><div></div><div>INSERT INTO `mem__quan`...'o') , `hostTo`) , ... (1)</div></div> | mem | | 6,903 | 0 | 0 | 0.79 | 18:16.538 | 17.784 | 0.159 | 6.699 | <div></div> | 7,356 | 0 |
| <div><div></div><div></div><div>INSERT INTO `mem__quan`...'on' = IF (VALUES (... (1)</div></div> | mem | | 6,985 | 0 | 0 | 0.86 | 10:54.856 | 8.940 | 0.094 | 8.301 | <div></div> | 7,332 | 0 |
| <div><div></div><div></div><div>INSERT INTO `mem__quan`...'en')) , `lastSeen`) (1)</div></div> | mem | | 7,025 | 37 | 0 | 1.00 | 3:11.791 | 11.220 | 0.027 | 4.436 | <div></div> | 13,947 | 0 |
| <div><div></div><div></div><div>INSERT INTO `mem__instr`...'tency') , `latency`) (1)</div></div> | mem | | 740 | 0 | 0 | 0.77 | 1:48.459 | 8.745 | 0.147 | 0.147 | <div></div> | 1,386 | 0 |
| <div><div></div><div></div><div>SELECT `mysqlconne0`asProc18_1191_0` , ... (1)</div></div> | mem | | 974 | 0 | 0 | 0.90 | 1:01.829 | 12.359 | 0.063 | 0.353 | <div></div> | 974 | 974 |
| <div><div></div><div></div><div>SELECT `mysqlserve0` `hostCache` AS ... (1)</div></div> | mem | | 1,801 | 0 | 0 | 0.93 | 53.661 | 13.351 | 0.030 | 0.636 | <div></div> | 1,801 | 1,801 |
| <div><div></div><div></div><div>INSERT INTO `mem__instr`...'es') , `diskWrites`) (1)</div></div> | mem | | 26 | 0 | 0 | 0.79 | 44.838 | 14.954 | 1.725 | 0.070 | <div></div> | 26 | 0 |
| <div><div></div><div></div><div>UPDATE `mem__inventory`...'n' = ? WHERE `hid` = ? (1)</div></div> | mem | | 321 | 0 | 0 | 0.80 | 43.886 | 9.668 | 0.137 | 0.044 | <div></div> | 321 | 321 |
| <div><div></div><div></div><div>UPDATE `mem__config` ? WHERE `user_id` = ? (1)</div></div> | mem | | 321 | 0 | 0 | 0.71 | 40.788 | 10.607 | 0.127 | 0.038 | <div></div> | 321 | 321 |
| <div><div></div><div></div><div>CREATE TEMPORARY TABLE ...(`id` INT8 NOT NULL) (1)</div></div> | mem | | 13 | 0 | 0 | 0.14 | 36.525 | 12.055 | 2.810 | 0.000 | <div></div> | 0 | 0 |
| <div><div></div><div></div><div>INSERT INTO `mem__instr`...'nedTableDefinitions`) (1)</div></div> | mem | | 26 | 0 | 0 | 0.64 | 34.348 | 13.349 | 1.321 | 0.003 | <div></div> | 27 | 0 |
| <div><div></div><div></div><div>UPDATE `mem__inventory`...'p' = ? WHERE `hid` = ? (1)</div></div> | mem | | 416 | 0 | 0 | 0.81 | 33.469 | 13.060 | 0.080 | 0.042 | <div></div> | 416 | 416 |
| <div><div></div><div></div><div>INSERT INTO `mem__instr`...'hed`) , `notCached`) (1)</div></div> | mem | | 26 | 0 | 0 | 0.61 | 32.509 | 11.782 | 1.250 | 0.003 | <div></div> | 27 | 0 |
| <div><div></div><div></div><div>SELECT * FROM (SELECT ...m_no_index_used` AS ... (1)</div></div> | mem | | 14 | 0 | 0 | 0.00 | 29.832 | 9.602 | 2.131 | 0.008 | <div></div> | 7,459 | 227,050 |
| <div><div></div><div></div><div>INSERT INTO `mem__instr`...'s`) , `connections`) (1)</div></div> | mem | | 25 | 0 | 0 | 0.82 | 29.462 | 14.294 | 1.178 | 0.005 | <div></div> | 26 | 0 |
| <div><div></div><div></div><div>INSERT INTO `mem__instr`...' (`sent`) , `sent`) (1)</div></div> | mem | | 25 | 0 | 0 | 0.82 | 28.991 | 14.332 | 1.160 | 0.003 | <div></div> | 26 | 0 |
| <div><div></div><div></div><div>SELECT `agent0` . `hid`.... `hid` = ? FOR UPDATE (1)</div></div> | mem | | 909 | 0 | 0 | 0.96 | 28.632 | 6.964 | 0.031 | 0.158 | <div></div> | 909 | 909 |
| <div><div></div><div></div><div>UPDATE `mem__events`me` = ? WHERE `id` = ? (1)</div></div> | mem | | 395 | 0 | 0 | 0.92 | 27.544 | 10.061 | 0.070 | 0.038 | <div></div> | 395 | 395 |

Copyright © 2005, 2013, Oracle and/or its affiliates. All rights reserved.

3.0.1.7150 - Cerberus.local (192.168.1.67) - Sep 16, 2013 11:50:28 am (Up Since: 38 minutes ago) - About

ORACLE

Copyright © 2015, Oracle and/or its affiliates. All rights reserved.

Query Analyzer Query Details

Canonical Query | **Example Query** | Explain Query | Graphs

The query with the longest execution time during the Time Span (usually the slowest but not always).

Sampled Query

[truncated](#) | [full](#) | [formatted](#)

SELECT

```
mysqlserve0.`hid` AS hid1124_0, mysqlserve0.`id` AS id2_1124_0,
mysqlserve0.`lastContact` AS lastCont3_1124_0,
mysqlserve0.`hasLastContact` AS hasLastC4_1124_0,
mysqlserve0.`startTime` AS startTime5_1124_0,
mysqlserve0.`hasStartTime` AS hasStart6_1124_0,
mysqlserve0.`timestamp` AS timestamp7_1124_0,
mysqlserve0.`capabilities` AS capabil18_1124_0,
mysqlserve0.`hasCapabilities` AS hasCapab9_1124_0,
mysqlserve0.`characterSet` AS charact10_1124_0,
mysqlserve0.`hasCharacterSet` AS hasChar11_1124_0,
mysqlserve0.`collation` AS collation12_1124_0,
mysqlserve0.`hasCollation` AS hasColl13_1124_0,
mysqlserve0.`connection` AS connection14_1124_0,
mysqlserve0.`hasConnection` AS hasConn15_1124_0,
mysqlserve0.`environment` AS environ16_1124_0,
```

Execution Time

27,084 ms

Date

Sep 16, 2013 1:07:17 PM

User

service_manager

Thread ID

10,712

From Host

localhost

To Host

Source Location

None found.

Comments

None found.


Canonical Query | Example Query | Explain Query | **Graphs**

Graphs of a query during the Time Span

Execution Time



Executions



Rows Affected



Performance Schema

Some useful tables

- **events_statements_history**
events_statements_history_long
 - Most recent statements executed
- **events_statements_summary_by_digest**
 - Summary for similar statements (same statement digest)
- **file_summary_by_event_name**
 - Interesting event: wait/io/file/innodb/innodb_data_file
- **table_io_waits_summary_by_table**
table_io_waits_summary_by_index_usage
 - Statistics on storage engine access per table and index

Performance Schema

Statement digest

- Normalization of queries to group statements that are similar to be grouped and summarized:

SELECT * FROM orders WHERE o_custkey=10 AND o_totalprice>20

SELECT * FROM orders WHERE o_custkey = 20 AND o_totalprice > 100

➡ SELECT * FROM orders WHERE o_custkey = ? AND o_totalprice > ?

- **events_statements_summary_by_digest**

DIGEST, DIGEST_TEXT, COUNT_STAR, SUM_TIMER_WAIT, MIN_TIMER_WAIT, AVG_TIMER_WAIT, MAX_TIMER_WAIT, SUM_LOCK_TIME, SUM_ERRORS, SUM_WARNINGS, SUM_ROWS_AFFECTED, SUM_ROWS_SENT, SUM_ROWS_EXAMINED, SUM_CREATED_TMP_DISK_TABLES, SUM_CREATED_TMP_TABLES, SUM_SELECT_FULL_JOIN, SUM_SELECT_FULL_RANGE_JOIN, SUM_SELECT_RANGE, SUM_SELECT_RANGE_CHECK, SUM_SELECT_SCAN, SUM_SORT_MERGE_PASSES, SUM_SORT_RANGE, SUM_SORT_ROWS, SUM_SORT_SCAN, SUM_NO_INDEX_USED, SUM_NO_GOOD_INDEX_USED, FIRST_SEEN, LAST_SEEN

Performance Schema

Statement events

- Tables:

events_statements_current (Current statement for each thread)

events_statements_history (10 most recent statements per thread)

events_statements_history_long (10000 most recent statements)

- Columns:

THREAD_ID, **EVENT_ID**, **END_EVENT_ID**, **EVENT_NAME**, **SOURCE**, **TIMER_START**, **TIMER_END**, **TIMER_WAIT**, **LOCK_TIME**, **SQL_TEXT**, **DIGEST**, **DIGEST_TEXT**, **CURRENT_SCHEMA**, **OBJECT_TYPE**, **OBJECT_SCHEMA**, **OBJECT_NAME**, **OBJECT_INSTANCE_BEGIN**, **MYSQL_ERRNO**, **RETURNED_SQLSTATE**, **MESSAGE_TEXT**, **ERRORS**, **WARNINGS**, **ROWS_AFFECTED**, **ROWS_SENT**, **ROWS_EXAMINED**, **CREATED_TMP_DISK_TABLES**, **CREATED_TMP_TABLES**, **SELECT_FULL_JOIN**, **SELECT_FULL_RANGE_JOIN**, **SELECT_RANGE**, **SELECT_RANGE_CHECK**, **SELECT_SCAN**, **SORT_MERGE_PASSES**, **SORT_RANGE**, **SORT_ROWS**, **SORT_SCAN**, **NO_INDEX_USED**, **NO_GOOD_INDEX_USED**, **NESTING_EVENT_ID**, **NESTING_EVENT_TYPE**

MySQL sys Schema / ps_helper

- Started as a collection of views, procedures and functions, designed to make reading raw Performance Schema data easier
- Implements many common DBA and Developer use cases
- Now bundled within MySQL Workbench
- Available on GitHub
 - <https://github.com/MarkLeith/mysql-sys>
- Examples of very useful functions:
 - `format_time()` , `format_bytes()`, `format_statement()`

MySQL sys Schema

Example

statement_analysis: Lists a normalized statement view with aggregated statistics, mimics the MySQL Enterprise Monitor Query Analysis view, ordered by the total execution time per normalized statement

```
mysql> select * from statement_analysis limit 1\G
```

```
***** 1. row *****
```

```
query: INSERT INTO `mem__quan`. `nor ... nDuration` = IF ( VALUES ( ...
```

```
db: mem
```

```
full_scan:
```

```
exec_count: 1110067
```

```
err_count: 0
```

```
warn_count: 0
```

```
total_latency: 1.93h
```

```
max_latency: 5.03 s
```

```
avg_latency: 6.27 ms
```

```
lock_latency: 00:18:29.18
```

```
rows_sent: 0
```

```
rows_sent_avg: 0
```

```
rows_examined: 0
```

```
rows_examined_avg: 0
```

```
tmp_tables: 0
```

```
tmp_disk_tables: 0
```

```
rows_sorted: 0
```

```
sort_merge_passes: 0
```

```
digest: d48316a218e95b1b8b72db5e6b177788!
```

```
first_seen: 2014-05-20 10:42:17
```

Optimizer Trace: Query Plan Debugging

- EXPLAIN shows the selected plan
- TRACE shows WHY the plan was selected:
 - Alternative plans
 - Estimated costs
 - Decisions made
- JSON format

Optimizer Trace: Example

```
SET optimizer_trace= "enabled=on", end_markers_in_json=on;
```

```
SELECT * FROM t1, t2 WHERE f1=1 AND f1=f2 AND f2>0;
```

```
SELECT trace INTO DUMPFILE <filename>
```

```
FROM information_schema.optimizer_trace;
```

```
SET optimizer_trace="enabled=off";
```

| | |
|-----------------------------------|--|
| QUERY | SELECT * FROM t1,t2 WHERE f1=1 AND f1=f2 AND f2>0; |
| TRACE | "steps": [{ "join_preparation": { "select#": 1,... } ... } ...] |
| MISSING_BYTES_BEYOND_MAX_MEM_SIZE | 0 |
| INSUFFICIENT_PRIVILEGES | 0 |

Program Agenda

- 1 Introduction to MySQL optimizer
- 2 Selecting data access method
- 3 Join optimizer
- 4 Sorting
- 5 Tools for monitoring, analyzing, and tuning queries
- 6 Influencing the Optimizer

Influencing the Optimizer

When the optimizer does not do what you want

- Add indexes
- Force use of specific indexes:
 - USE INDEX, FORCE INDEX, IGNORE INDEX
- Force specific join order:
 - STRAIGHT_JOIN
- Adjust session variables
 - optimizer_switch flags: set optimizer_switch="index_merge=off"
 - Buffer sizes: set sort_buffer=8*1024*1024;
 - Other variables: set optimizer_prune_level = 0;

More information

- My blog:
 - <http://oysteing.blogspot.com/>
- Optimizer team blog:
 - <http://mysqloptimizerteam.blogspot.com/>
- MySQL Server Team blog
 - <http://mysqlserverteam.com/>
- MySQL forums:
 - Optimizer & Parser: <http://forums.mysql.com/list.php?115>
 - Performance: <http://forums.mysql.com/list.php?24>

Safe Harbor Statement

The preceding is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

Q&A

ORACLE®

Copyright © 2015, Oracle and/or its affiliates. All rights reserved.

